```python
from transformers import AutoTokenizer, AutoModelForCausalLM
model_id = "meta-llama/Meta-Llama-3.1-8B"
tokenizer = AutoTokenizer.from_pretrained(model_id)
lm = AutoModelForCausalLM.from_pretrained(model_id)

mlp = lm.model.layers[16].mlp.down_proj
neurons = [394, 5490, 8929]
prompt = "The truth is the"


def pre_hook_fn(module, input):
    input[0][:,-1,neurons] = 10

hook = mlp.register_forward_pre_hook(pre_hook_fn)
inputs = tokenizer(prompt, return_tensors="pt")
out = lm(**inputs)
hook.remove()

last = out["logits"][:, -1].argmax()
prediction = tokenizer.decode(last)
print(prediction)
```

(a)

```python
from nnsight import LanguageModel
model_id = "meta-llama/Meta-Llama-3.1-8B"
lm = LanguageModel(model_id)


mlp = lm.model.layers[16].mlp.down_proj
neurons = [394, 5490, 8929]
prompt = "The truth is the"




with lm.trace(prompt, remote=True):
    mlp.input[:, -1, neurons] = 10
    out = lm.output.save()


last = out["logits"][:, -1].argmax()
prediction = lm.tokenizer.decode(last)
print(prediction)
```

(b)